

## EECE 5645 HOMEWORK 2 Solutions

Question 1 : Refer the pdf **hw2question1.pdf** attached if you don't get my 1<sup>st</sup> question below

Part A)

Given a vector  $x \in \mathbb{R}^d$  and a real number  $y \in \mathbb{R}$ , define the function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  as

$$f(\beta; x, y) = \left( y - x^T \beta \right)^2 = \left( y - \sum_{k=1}^d \beta_k x_k \right)^2.$$

To find the partial derivatives of  $f$  with respect to  $\beta_k$ , so expand the square term in the numerator

numerator  
so expand

$$\begin{aligned} & \left( y - \sum_{k=1}^d \beta_k x_k \right)^2 \\ &= y^2 - \left( \sum_{k=1}^d \beta_k x_k \right)^2 \\ &= (y - x^T \beta)^2 \end{aligned}$$

It can be written as

can be written as

$$\begin{aligned} &= (y - b_1 x_1 - b_2 x_2 - \dots - b_d x_d)^2 \\ &= y^2 - 2y x^T \beta + (x^T \beta)^2 \\ &= y^2 - 2y x^T \beta + (x^T \beta)(\beta^T x) \\ f(\beta; x, y) &= y^2 - 2y x^T \beta + x^T (\beta \beta^T) x \\ &= y^2 - 2x^T y \beta + (x^T)^2 (\beta^T)^2 \end{aligned}$$

Hence  $\beta^T$  is a scalar quantity and  $\beta_k \beta_k^T$  is a  $d \times d$  matrix so it can be written as  $(\beta_k)^2$

Hence by taking partial derivatives

$$\frac{df}{d\beta_k} = -2y x_k + 2 \beta_k x_k^2$$

also  $\beta_k^T \beta_k$  can be expanded as  
 $= (b_1^2 b_1 b_2 \dots b_1 b_d, b_2^2 b_1 b_2 \dots b_2 b_d, \dots b_d^2 b_1 b_2 \dots b_d b_d)$

hence the

$$\frac{df}{d\beta_k} = -2y x_k + 2 \beta_k x_k^2 \quad \text{for } k = 1, \dots, d$$

where  $\beta_k \beta_k^T$  denotes the  $k^{\text{th}}$  component of the vector  $\beta_k \beta_k^T x$ . Therefore the partial derivatives of  $f$  with respect to  $\beta_k$  is  $-2$  times the  $k^{\text{th}}$  component of  $x$  times  $y$  plus  $2$  times  $k^{\text{th}}$  component of vector  $\beta_k x_k$ .

Ans B )

Question 2 : The gradient of  $f$  as a  $d$ -dimensional column vector with components given by the partial derivatives with respect to each element of  $\beta$ .

$$\text{so from } \frac{df}{d\beta_k} = -2y x_k + 2 \beta_k x_k^2$$

$\text{sf}(B, x, y) = [df|dB_1, df|B_2 \dots df|B_d]$   
using dot product of  $B$  and  $x$  as

$$x^T B = [x_1, x_2, \dots, x_d] [B_1, B_2, \dots, B_d]$$

taking derivatives of the expression we get  
w.r.t  $B_k$

$$\frac{d}{dB_k} (x^T B) = x_k$$

$$\begin{aligned} \frac{df}{dB_k} &= -2y x_k + (2x^T B) x_k \\ &= 2(y - x^T B) x_k \end{aligned}$$

$$\text{hence } \text{sf}(B, x, y) = [-2(y - x^T B)x_1, -2(y - x^T B)x_2, \dots, -2(y - x^T B)x_d]^T$$

to solve it further

$$\text{sf}(B, x, y) = -2(y - x^T B) [x_1, x_2, \dots, x_d]^T$$

using matrix multiplication we get

$$\text{sf}(B, x, y) = -2(y - x^T B)x$$

using the transpose of a column vector is a row vector

$$\begin{aligned} \text{sf}(B, x, y) &= -2(y - x^T B) x^T \\ \therefore \text{The gradient descent can be} \\ &\text{given as.} \end{aligned}$$

$$= -2(y x_k - B_k^T x_k)$$

$$= -2(y - x^T B_k) x_k //$$

Part C)

Ans c Given:

$$F(\beta) = \frac{1}{n} \sum_{i=1}^n f(\beta; x_i, y_i) + \lambda \|\beta\|_2^2$$
$$= \frac{1}{n} \sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \beta^T \beta$$

Here  $\lambda \geq 0$  and  $\beta \in \mathbb{R}^d$ ,

To prove  $\nabla F(\beta) = -\frac{2}{n} \sum_{i=1}^n (y_i - x_i^T \beta) x_i + 2\lambda \beta$

hence when we take partial derivatives,

$$\frac{dF}{d\beta} = \frac{1}{n} \sum_{i=1}^n \frac{df}{d\beta} + 2\lambda \beta$$

$$= \left(\frac{1}{n}\right) \sum_{i=1}^n \left[ 2(x_i^T \beta) (y_i - x_i^T \beta) \right] + 2\lambda \beta$$

$$= \frac{-2}{n} \sum_{i=1}^n (x_i^T (y - x_i^T \beta)) + 2\lambda \beta.$$

hence

$$\nabla F(\beta) = \left( \partial F / \partial \beta_1, \partial F / \partial \beta_2, \dots, \partial F / \partial \beta_d \right)$$

$d$  dimension of vector  $x_i$  and  $\beta$ . So partial derivatives can be given as

$$\partial F / \partial \beta_k = \partial / \partial \beta_k \left( \frac{1}{n} \sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \|\beta\|^2 \right)$$

$$\sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \|\beta\|^2$$

$$= \frac{1}{n} \sum_{i=1}^n \partial / \partial \beta_k (y_i - x_i^T \beta)^2 + \partial / \partial \beta_k \lambda \|\beta\|^2$$

$$= \frac{1}{n} \sum_{i=1}^n 2 (y_i - x_i^T \beta) (-x_{ki}) + 2 \lambda \beta_k$$

using chain rule  $\|\beta\|^2 = \beta^T \beta$

$$= -\frac{2}{n} \sum_{i=1}^n (y_i - x_i^T \beta) x_{ki} + 2 \lambda \beta_k$$

$$\nabla F(\beta) = (-2/n) \sum_{i=1}^n (y_i - x_i^T \beta) x_i + 2 \lambda \beta$$

which is desired result.

Part D)

Given that

$$h(\gamma) = F(\beta_1 + \gamma \beta_2),$$

$$h(\gamma) = a\gamma^2 + b\gamma + c,$$

Substituting  $F(\beta_1 + \gamma \beta_2)$  with its definition we get

$$h(\gamma) = \frac{1}{n} \sum_{i=1}^n (y_i - x_i^T (\beta_1 + \gamma \beta_2))^2 + \lambda (\beta_1 + \gamma \beta_2)^2$$

Expanding the squared terms we get is



$$h(\gamma) = \frac{1}{n} \sum (y - x^T \beta_1 - \gamma x_i^T \beta_2)^2 + 1 \|\beta_1\|^2 + 2 \gamma \beta_1^T \beta_2 + 1 \gamma^2 \|\beta_2\|^2$$

On expanding further we get

$$h(\gamma) = \frac{1}{n} \sum (y_i^2 - 2 y_i x_i^T \beta_1 + (x_i^T \beta_1)^2 + 2 \gamma x_i^T \beta_2 y_i -$$

$$(2 \gamma x_i^T \beta_1 y_i) + (x_i^T \beta_2)^2 + \gamma^2 x_i^T \beta_2)^2 + 1 \|\beta_1\|^2 + 2 \gamma \beta_1^T \beta_2 + 1 \gamma^2 \|\beta_2\|^2$$

On simplifying we get

$$h(\gamma) = \frac{1}{n} \sum_{i=1}^n (x_i^T \beta_1)^2 - 2 y_i x_i^T \beta_1 + y_i^2 + 2 \gamma x_i^T \beta_2 y_i - 2 \gamma (x_i^T \beta_1) (x_i^T \beta_2) + \gamma^2 (x_i^T \beta_2)^2 + 1 \|\beta_1\|^2 + 2 \gamma \beta_1^T \beta_2 + 1 \gamma^2 \|\beta_2\|^2$$

Collecting the terms having same power we get

$$\begin{aligned}
 h(\gamma) = & \frac{1}{n} \sum_{i=1}^n (\alpha_i^T \beta_1)^2 - 2 y_i \alpha_i^T \beta_1 + y_i^2 + \\
 & \lambda \|\beta_1\|^2 + 2 \alpha_i^T \beta_2 y_i - 2 (\alpha_i^T \beta_1) (\alpha_i^T \beta_2) \\
 & + 2 \lambda \beta_1^T \beta_2 \gamma + \gamma^2 \sum (\alpha_i^T \beta_2)^2 + \lambda \gamma^2 \|\beta_2\|^2
 \end{aligned}$$

Comparing it in general from we get

$$h(\gamma) = a\gamma^2 + b\gamma + c$$

we can see that

$$a = \frac{1}{n} \left[ \sum_{i=1}^n (\alpha_i^T \beta_2)^2 + \lambda \beta_2^T \beta_2 \right]$$

$$\begin{aligned}
 b = & \frac{-2}{n} \sum_{i=1}^n (\alpha_i^T \beta_2) (y_i - \alpha_i^T \beta_1) \\
 & + 2 \lambda \beta_1^T \beta_2
 \end{aligned}$$

$$c = \frac{1}{n} \sum_{i=1}^n (y_i - \alpha_i^T \beta_1)^2 + \lambda \beta_1^T \beta_1$$

Ans E)

Ans E To show  $a > 0$

Here we need to show coefficient of  $\gamma^2$   
 $a$  is +ve

since  $(x_1^T \beta_2)^2$  and  $\frac{1}{\|\beta_2\|^2}$  are both  
non-negative since  $1 > 0$ , their sum  
is also non negative.  
 $\therefore a > 0$  as required.

Ans f :

As per the convex function definition we need to show that 2<sup>nd</sup> derivative for all values of gamma . So taking derivative of gamma we get .

$$h'(\gamma) = 2a\gamma + b$$

2<sup>nd</sup> derivative

$$h''(\gamma) = 2a$$

Since  $a$  is  $a > 0$  and is a non negative so  $h''(\gamma)$  is also non negative for all value of gamma . Hence  $h$  convex function when  $a$  is not equal to 0



TO find  $\min(h)$ ,  $\frac{dh}{dx} \approx 0$

$$h'(x^*) = 2ax + b = 0$$

$x^* = -b/2a$  when we solve for  $x^*$   
since  $h$  is a convex function, this is  
a critical point and is minimum.  
 $\therefore h$  is minimized at  $x^* = -b/2a$ .

Ans G)

Given that

Minimize:  $h(\gamma)$

subject to:  $\gamma \in [0, 1]$ .

Assume that  $a \neq 0$ . Find the optimal  $\gamma \in [0, 1]$  in each of these cases:

- (a)  $\gamma^* \in [0, 1]$ ,
- (b)  $\gamma^* < 0$ ,
- (c)  $\gamma^* > 1$ ,

(a)  $x \in [0, 1]$

$\frac{-b}{a} \in [0, 1]$

minimum of  $h$  is achieved at

$x^* = \frac{-b}{2a}$

since it is already in a feasible range

(b)

$\frac{-b}{2a} < 0$

, then minimum  $h$  is achieved at

$x^* = 0$

since it is smallest range value within the feasible range.

Part b

If  $\gamma^*$  is less than 0, then the optimal solution is

$\gamma^* = 0$ , since  $\gamma$  cannot be less than 0.

Part c If  $\gamma^*$  is greater than 1, then the optimal solution is  $\gamma^* = 1$ , since  $\gamma$  cannot be greater than 1

(c)  $\frac{-b}{2a} > 1$ , min  $h$  is achieved at  $y^* = 1$   
 since it is the largest value in the feasible range.

when  $a=0$  we can see optimal problem such as optimization becomes linear function over a bounded interval  $[0,1]$  and the min is achieved at one of the endpoints.

If  $b \geq 0$  the  $\min(h(y))$  is at  $y^* = 0$   
 If  $b \leq 0$  the  $\min(h)$  is at  $y^* = 1$

then any value in the interval  $[0,1]$  is optimal.

Ans given  $|z_{k^*}| \geq |z_k|$

$$\begin{aligned} & \min z^T \beta \\ & \text{subject to } \sum_{k=1}^d |\beta_k| \leq K \end{aligned}$$

Let (a) To prove if  $\beta$  is feasible.

$$|\beta_k| = |-k \operatorname{sign}(z_{k^*})| = k \text{ and}$$

$$|B_k| = |0| = 0 \text{ for all } k \neq k^*$$

so we have

$$\sum_{k=1}^d |B_{k^*}| = |B_{k^*}| + \sum_{k=1, (k \neq k^*)}^d |B_k|$$

$$= K + 0$$

$$= K$$

Hence,  $\beta^*$  satisfies the constraints

$$\sum_{k=1}^d |B_k| \leq K \text{ and is feasible}$$

Part b proving optimality of  $\beta^*$ :

let  $\beta$  be any feasible solution satisfying the constraint

$$\sum_{k=1}^d |B_k| \leq K$$

$$\text{so } z^T \beta \geq z^T \beta^*$$

let  $k^*$  be coordinate

$$z^T \beta^* = z_{k^*} (-K \text{ sign}(z_{k^*})) = -K/|z_{k^*}|$$



$$z^T \beta = \sum_{k=1}^d z_k \beta_k \leq \sum_{k=1}^d |z_k| |\beta_k|$$

As  $|z_{k^*}| \geq |z_k|$ , we have

$$|z_{k^*}| |\beta_k| \geq |z_k| |\beta_k|, \text{ and}$$

$$|z_{k^*}| |\beta_k| \geq \sum_{k=1}^d |z_k| |\beta_k|$$

$\therefore$

$$|z_{k^*}| |\beta_k| + \sum_{k=1}^d |z_k| |\beta_k| \geq$$

$$\sum_{k=1}^d |z_k| |\beta_k| \geq z^T \beta.$$

multiplying both side  $-k \operatorname{sign}(z_{k^*})$

$$\beta_k - \sum_{k=1}^d k \operatorname{sign}(z_k) |z_k| \beta_k \geq -k z^T \beta$$

since  $\beta_k = 0$  for all  $k \neq k^*$ , the sum reduce to:

$$-k |z_{k^*}| \operatorname{sign}(z_{k^*}) \beta_k \geq -k z^T \beta$$

divide both sides by  $-k$  and  $\times$  by  $z_{k^*}$

$$|z_k^*| (\text{sign } z_k^*) \beta_k \leq z^T \beta$$

$\text{sign } z_k^*$  and  $\text{sign}(z_k)$

are both either +1 or -1,

so when we multiply by  $|z_k^*|$  to get

$$z_k^* \beta_k \leq z_k \beta_k$$

since  $|z_k^*| > |z_k|$

$$|z_k^*| |\beta_k| > |z_k| |\beta_k| > \sum_{k=1}^d |z_k| |\beta_k|$$

$\rightarrow \textcircled{1}$

Then

$$\textcircled{1} > z^T \beta$$

Hence we can say that

$z^T \beta^* = -K |z_k^*| \leq z^T \beta$   
for any feasible  $\beta$ , proving the optimality of  $\beta^*$

## Question 2 :

After reading Help(PR.readData) it prints this :

Help on function readData in module ParallelRegression:

```
readData(input_file, spark_context)
    Read data from an input file and return rdd containing pairs of the form:
        (x,y)
    where x is a numpy array and y is a real value. The input file should be a
    'comma separated values' (csv) file: each line of the file should contain x
    followed by y. For example, line:
```

```
1.0,2.1,3.1,4.5
```

should be converted to tuple:

After reading Help(PR.F) it prints this :

Help on function f in module ParallelRegression:

```
f(x, y, beta)
    Given vector x containing features, true label y,
    and parameter vector  $\beta$ , return the square error:
```

$$f(\beta; x, y) = (y - \langle x, \beta \rangle)^2$$

First the argument for the whole program in spark is defined as below

```
args = parser.parse_args()

sc = SparkContext(appName='Parallel Regression')
sc.setLogLevel('warn')
```

Then there are various options that are stored within the argument. It describes the format on which the input needs to be given in terms of choosing an option for , input files , options to choose like test data or train data . Hence it describes the use cases of the positional argument.

The part of the code that helps in reading and printing the above commands is highlighted as below:

```
if __name__ == "__main__":
    parser = argparse.ArgumentParser(description = 'Parallel Regression.', formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    parser.add_argument('--traindata', default=None, help='Input file containing (x,y) pairs, used to train a linear model')
    parser.add_argument('--testdata', default=None, help='Input file containing (x,y) pairs, used to test a linear model')
```

This helps us in selecting the right option from the menu and then it goes within the function as below :

```
if args.testdata is not None:
    # Read beta from args.beta, and evaluate its MSE over data
    print('Reading test data from', args.testdata)
    data = readData(args.testdata, sc)
```

Once the values goes to the respective functions then in each function whatever are the comments written within it for us to help in coding and to make us understand what code does , these comments are the printed when we call it using `help(pr.readdata)` or `help(pr.r)`. This is shown below :

```
def readData(input_file,spark_context):
    """ Read data from an input file and return rdd containing pairs of the form:
        (x,y)
        where x is a numpy array and y is a real value. The input file should be a
        'comma separated values' (csv) file: each line of the file should contain x
        followed by y. For example, line:

        1.0,2.1,3.1,4.5

        should be converted to tuple:

        (array(1.0,2.1,3.1),4.5)
    """

def f(x,y,beta):
    """ Given vector x containing features, true label y,
        and parameter vector  $\beta$ , return the square error:

        
$$f(\beta;x,y) = (y - \langle x,\beta \rangle)^2$$

    """
```

After the values are passed inside the function it helps us to generate the data for the `f` and When you store it in a another rdd and take inputs you get output of the rdd as below :

The format of a line in the “data/small.test “ file sis expected to be a comma- separated list of numerical values representing the independent variables , followed by a single numerical value representing the dependent variables . For example below is the screenshot attached of the file :

```
-1.336451934422684173e+00,3.600848542167133792e-01,-1.516961555037868170e-01,1.786103773022134522e+00,-4.8123
2.106503876368362338e+00,-4.329014446182832265e-01,9.442626357254089164e-01,4.437358581154937198e+00,-9.11908
-3.245733207805587828e-01,-7.704172635517650969e-01,1.193850392195123122e+00,1.053478405625195152e-01,2.50056
5.438613463532162573e-01,-3.034140805865944701e-01,1.984081610357717684e+00,2.957851640571330654e-01,-1.65015
1.609649379872065411e-01,-1.066439888547795034e+00,-7.392464836862415734e-01,2.590971126122524823e-02,-1.7165
-2.757574043496156535e-01,-2.513288265715642122e+00,-9.158789076489609604e-01,7.604214605363743273e-02,6.9305
-6.091680874067831875e-01,-1.088233294499143938e-01,4.632135168995548113e-01,3.710857587148382319e-01,6.62916
-1.490040323249569587e+00,2.853970370486333064e+00,1.684770351679401035e+00,2.220220164909681770e+00,-4.25253
-2.382511830481278459e+00,-1.765204666428647484e-01,8.257640118836888643e-01,5.676362622383251733e+00,4.20562
-5.010526314787675517e-01,-5.673718789863278072e-02,-1.787057719772146225e+00,2.510537395117976578e-01,2.8428
```

The format of an element in the resulting RDD is a tuple of two elements , where the first element is a list of numerical values representing the independent variables , and the second element is a single numerical value representing the dependent variable . When you store it in a another rdd and take inputs you get output of the rdd as below :





Output to reload the program using the functions in the pyspark . This a library obtained from the spark documentation Whenever you change the contents of ParallelRegression.py, you can update the function definitions in pyspark by using it as below

```
>>> import importlib
>>> import ParallelRegression as PR
>>> importlib.reload(PR)
<module 'ParallelRegression' from '/home/agnihotri.ak/EECE5645/ParallelRegression.py'>
```

#### Question 4 :

Answer : The code for function f

Part A )

```
def f(x,y,beta):
    """ Given vector x containing features, true label y,
        and parameter vector  $\beta$ , return the square error:
        
$$f(\beta;x,y) = (y - \langle x,\beta \rangle)^2$$

    """
    return (y - predict(x, beta)) ** 2
```

Part B) The code for local gradient

```
def localGradient(x,y,beta):
    """ Given vector x containing features, true label y,
        and parameter vector  $\beta$ , return the gradient  $\nabla f$  of f:
        
$$\nabla f(\beta;x,y) = -2 * (y - \langle x,\beta \rangle) * x$$

        with respect to parameter vector  $\beta$ .
        The return value is  $\nabla f$ .
    """
    return -2 * (y - np.dot(x, beta)) * x
```

Part C)

Testing local\_gradient and estimated\_gradient . Tested with at least 100 random values of x,y, and  $\beta$  (using, e.g., a for loop), and appropriate assert statements when  $\delta$ (delta) value is extremely small .

First the rand vector function generates the random values we need for test as we need 100 random values

```
def rand_vector():  
    x1 = np.random.rand(100)  
    beta = np.random.rand(100)  
    y1 = random()  
    return x1, y1, beta
```

The function to test :

```
def gradient_test(rddout):  
    d = 0.0001 # this is delta  
    x1, y1, beta = rand_vector()  
    gradvalue = PR.estimateGrad(lambda beta: PR.f(x1,y1,beta),beta, d)  
    localgradvalue = PR.localGradient(x1, y1, beta)  
    print("Local Gradient = ", localgradvalue)  
    print("Estimated Gradient = ", gradvalue)  
    for i in range(len(gradvalue)):  
        sub= abs(localgradvalue[i] - gradvalue[i])  
        print("Difference = ", sub)  
        assert sub < d, f"Assert failed"
```

The main function to call the test

```
if __name__ == "__main__":  
    sc = SparkContext(appName='test code')  
    sc.setLogLevel('warn')  
    input="data/small.test"  
    rddout=PR.readData(input,sc)  
    print(rddout.take(10))  
    gradient_test(rddout)  
    print("All tests passed 1!")
```

Output of the test that tell it has passed the cases well :

```
Difference = 6.018818527309122e-05  
Difference = 1.8257841173152656e-05  
Difference = 9.816530379680444e-06  
Difference = 4.1149819615782235e-05  
Difference = 6.366398400459161e-06  
Difference = 5.994151408117432e-05  
Difference = 8.75310360157755e-05  
Difference = 1.7132199712222018e-06  
Difference = 6.055665768833762e-07  
Difference = 8.12521435911151e-06  
Difference = 3.9173080024212936e-07  
Difference = 1.9874672865682896e-05  
All tests passed 1!
```

## Question 5 :

### Part A) Function F code

```
def F(data,beta,lam = 0):
    """ Compute the regularized mean square error:

        
$$F(\beta) = \frac{1}{n} \sum_{(x,y) \in \text{data}} f(\beta; x, y) + \lambda ||\beta||_2^2$$

        
$$= \frac{1}{n} \sum_{(x,y) \in \text{data}} (y - \langle x, \beta \rangle)^2 + \lambda ||\beta||_2^2$$


        where n is the number of (x,y) pairs in RDD data.

        Inputs are:
        - data: an RDD containing pairs of the form (x,y)
        - beta: vector  $\beta$ 
        - lam: the regularization parameter  $\lambda$ 

        The return value is F( $\beta$ ).
    """

    n = data.count()
    mean_squared_error = data.map(
        lambda element: f(element[0], element[1], beta) / n).sum()
    regularization_term = lam * np.dot(beta, beta)
    return mean_squared_error + regularization_term
```

### Part B) The gradient code

```
def gradient(data,beta,lam = 0):
    """ Compute the gradient  $\nabla F$  of the regularized mean square error

        
$$F(\beta) = \frac{1}{n} \sum_{(x,y) \in \text{data}} f(\beta; x, y) + \lambda ||\beta||_2^2$$

        
$$= \frac{1}{n} \sum_{(x,y) \in \text{data}} (y - \langle x, \beta \rangle)^2 + \lambda ||\beta||_2^2$$


        where n is the number of (x,y) pairs in data.

        Inputs are:
        - data: an RDD containing pairs of the form (x,y)
        - beta: vector  $\beta$ 
        - lam: the regularization parameter  $\lambda$ 

        The return value is an array containing  $\nabla F$ .
    """

    n = data.count()
    mse_gradient = data.map(
        lambda element: localGradient(element[0], element[1], beta) / n
    ).sum()
    weights_gradient = 2 * lam * beta
    return mse_gradient + weights_gradient
```

### Part C)

Testing gradient actual value and gradient calculated value . Tested with multiple random values of  $\beta$  and  $\lambda$  (using, e.g., a for loop), and appropriate assert statements. Using 'data/small.test 'as a dataset. This defined in the function below

```
def gradient_file_test(rddout):
    d = 0.0001 # this is delta
    beta = np.random.rand(9)
    lam = random() * 5.0 + 0.1
    gradientactualvalue = PR.gradient(rddout, beta, lam)
    gradientcalculatedvalue = PR.estimateGrad(lambda beta: PR.F(rddout, beta, lam), beta, 0.0000001)
    print("Actual Gradient = ", gradientactualvalue)
    print("Estimated Gradient = ", gradientcalculatedvalue)
    for i in range(len(gradientactualvalue)):
        sub = abs(gradientcalculatedvalue[i] - gradientactualvalue[i])
        print("Difference = ", sub)
        assert sub < d, f"Assert failed"
```

The main function to test the gradient\_file\_test function()

```
if __name__ == "__main__":
    sc = SparkContext(appName='test code')
    sc.setLogLevel('warn')
    input="data/small.test"
    rddout=PR.readData(input,sc)
    print(rddout.take(10))
    #gradient_test(rddout)
    #print("All tests passed 1!")
    gradient_file_test(rddout)
    print("All tests passed 2!")
```

Output : It's the output for the above test it says that all test cases are passed . Also an rdd of 10 is printed to verify well

```
Actual Gradient = [ -4.58383292 -45.33856758 -14.07694946 13.27115949 32.86146581
15.93930387 -2.7398814 -17.97736514 21.83277825]
Estimated Gradient = [ -4.5838334 -45.33856782 -14.07694981 13.2711591 32.86146523
15.93930278 -2.73988121 -17.97736445 21.83277843]
Difference = 4.848518111444378e-07
Difference = 2.3960751605045516e-07
Difference = 3.4630211054320625e-07
Difference = 3.8445504202400116e-07
Difference = 5.77279145375087e-07
Difference = 1.0810062214261507e-06
Difference = 1.8549331182171613e-07
Difference = 6.90951150517094e-07
Difference = 1.843307373405878e-07
All tests passed 2!
```

### Question 6 :

Using the equations below

$$a = \lambda \beta_2^T \beta_2 + 1/n \sum (X_i^T \beta_2)^2$$

$$b = 2\lambda \beta_1^T \beta_2 - 2/n \sum (X_i^T \beta_2) (Y_i - X_i^T \beta_1)$$

$$c = 1/n \sum (Y_i - X_i^T \beta_1)^2 + \lambda \beta_1^T \beta_1$$

for the hcoeff functions

```

def hcoeff(data,beta1, beta2, lam = 0):
    """ Compute the coefficients a,b,c of quadratic function h, defined as :
        
$$h(y) = F(\beta_1 + y\beta_2) = ay^2 + by + c$$

        where F is the regularized mean square error function.

        Inputs are:
        - data: an RDD containing pairs of the form (x,y)
        - beta1: vector  $\beta_1$ 
        - beta2: vector  $\beta_2$ 
        - lam: the regularization parameter  $\lambda$ 

        The return value is a tuple containing (a,b,c).
    """
    n = data.count()
    a1 = lam * np.dot(np.transpose(beta2),beta2)
    a2 = data.map(lambda xy: pow(predict(xy[0],beta2),2))\
              .sum()
    a = a1 + (a2 / n)
    b1 = 2 * lam * np.dot(np.transpose(beta1),beta2)
    b2 = data.map(lambda xy: predict(xy[0],beta2) * (xy[1] - predict(xy[0],beta1)))\
              .sum()
    b = b1 - ((2 / n) * b2)
    c1 = lam * np.dot(np.transpose(beta1),beta1)
    c2 = data.map(lambda x: f(x[0],x[1],beta1))\
              .sum()
    c = c1 + (c2 / n)
    return (a,b,c)

```

Testing the coeff well :

```

def test_hcoeff(rddout):
    beta1=np.random.rand(9)
    beta2=np.random.rand(9)
    gamma = np.random.randn()
    tol = 1e-1
    lamda = 0.01
    a, b, c = PR.hcoeff(rddout,beta1,beta2,lamda)
    f1 = PR.F(rddout, beta1 + gamma*beta2, lamda)
    f2 = a*((gamma)**2)+ b*(gamma) + c
    print("f1",f1)
    print("f2",f2)
    assert abs(f1 - f2) < tol ,f"Assert failed"

```

The main function to call the test test\_coeff()

```

if __name__ == "__main__":
    sc = SparkContext(appName='test code')
    sc.setLogLevel('warn')
    input="data/small.test"
    rddout=PR.readData(input,sc)
    print(rddout.take(10))
    #gradient_test(rddout)
    #print("All tests passed 1!")
    #gradient_file_test(rddout)
    #print("All tests passed 2!")
    test_hcoeff(rddout)
    print("All the test passed 3!")

```

Output : The output tells us that the value of f1 and f2 matches with each other . Also all the test cases are passed . Hence the equation written for the l2 norm regularization are correct .



f1 430.8098672273337

f2 430.8098672273337

All the test passed 3!

### Question 7 :

#### Part A : The function for the test

```
def test(data,beta):
    """ Compute the mean square error

        | MSE( $\beta$ ) =  $1/n \sum_{(x,y) \text{ in data}} (y - \langle x, \beta \rangle)^2$ 

        of parameter vector  $\beta$  over the dataset contained in RDD data, where n is the size of RDD data.

        Inputs are:
        | - data: an RDD containing pairs of the form (x,y)
        | - beta: vector  $\beta$ 

        The return value is MSE( $\beta$ ).
    """
    n = data.count()
    mse = data.map(lambda element: f(element[0], element[1], beta) / n).sum()
    return mse
```

#### Part B: The function for exactLineSearch referring the 1(g) question logic

```
def exactLineSearch(data,beta,g,lam = 0):
    """ Given data, a vector x, and a direction g, return
        |  $\gamma^*$  = argmin_{ $\gamma$ } F(data,  $\beta - \gamma g$ )

        The solution is found by first computing the coefficients of the quadratic
        polynomial
        |  $h(\gamma) = F(\text{data}, \beta - \gamma g) = a\gamma^2 + b\gamma + c$ 
        The return value is  $\gamma^* = -b/(2*a)$ 

        Inputs are:
        | - data: an RDD containing pairs of the form (x,y)
        | - beta: vector  $\beta$ 
        | - g: direction vector g
        | - lam: the regularization parameter  $\lambda$ 

        The return value is  $\gamma^*$ 
    """
    # Compute the coefficient of the quadratic function for gamma
    # here cof is the coefficient stores tuples cof=(a,b,c) obtained from hcoeff function
    cof = hcoeff(data, beta, -g, lam)
    if cof[0] == 0:
        gamma = -cof[2]/cof[1]
    else:
        gamma = -cof[1] / (2 * cof[0]) # this nothing but gamma = -b/2a
    return gamma
```

#### Part C:

Output for small test and train

```

Reading training data from data/small.train
Gradient descent training on data from data/small.train with  $\lambda = 10.0$ ,  $\epsilon = 0.01$ , max iter = 100
k = 0   t = 1.1679270267486572   F( $\beta_k$ ) = 475.80205650084224   || $\nabla F(\beta_k)$ ||_2= 48.36642222745724   y_k = 0.034644340508364575
k = 1   t = 2.241664409637451   F( $\beta_k$ ) = 435.28011656165035   || $\nabla F(\beta_k)$ ||_2= 6.369684533296867   y_k = 0.041392870179372175
k = 2   t = 3.2633073329925537   F( $\beta_k$ ) = 434.4404025625204   || $\nabla F(\beta_k)$ ||_2= 1.0347072406178048   y_k = 0.03515640336219431
k = 3   t = 4.279298543930054   F( $\beta_k$ ) = 434.42158300451774   || $\nabla F(\beta_k)$ ||_2= 0.15479527474513602   y_k = 0.04185252366896443
k = 4   t = 5.280001163482666   F( $\beta_k$ ) = 434.42108157828164   || $\nabla F(\beta_k)$ ||_2= 0.02789933943570143   y_k = 0.035332834868740795
k = 5   t = 6.218413352966309   F( $\beta_k$ ) = 434.42106782721686   || $\nabla F(\beta_k)$ ||_2= 0.004340132957768057   y_k = 0.041958276278279694
Algorithm ran for 6 iterations. Converged: True Training time: 6.218591928482056
Saving trained  $\beta$  in beta_small_gd
Reading test data from data/small.test
Reading  $\beta$  from beta_small_gd
Computing MSE on data data/small.test
MSE is: 399.47612435238216

```

## Part D :

### Output for large test and train

```

Reading training data from data/large.train
Gradient descent training on data from data/large.train with  $\lambda = 10.0$ ,  $\epsilon = 0.01$ , max iter = 100
k = 0   t = 1.2748327255249023   F( $\beta_k$ ) = 449.5326978248175   || $\nabla F(\beta_k)$ ||_2= 34.179638533361754   y_k = 0.03983855923496276
k = 1   t = 2.375239372253418   F( $\beta_k$ ) = 426.2620454198274   || $\nabla F(\beta_k)$ ||_2= 10.024501896219903   y_k = 0.02377579167421469
k = 2   t = 3.4401891231536865   F( $\beta_k$ ) = 425.06742317950113   || $\nabla F(\beta_k)$ ||_2= 1.8399312668497296   y_k = 0.040062729653947055
k = 3   t = 4.49951434135437   F( $\beta_k$ ) = 424.9996100573415   || $\nabla F(\beta_k)$ ||_2= 0.5704360138258547   y_k = 0.023829714168459224
k = 4   t = 5.536103248596191   F( $\beta_k$ ) = 424.9957329956614   || $\nabla F(\beta_k)$ ||_2= 0.1068094270498083   y_k = 0.040179708024953915
k = 5   t = 6.5032265186309814   F( $\beta_k$ ) = 424.99550380550977   || $\nabla F(\beta_k)$ ||_2= 0.033774449038017844   y_k = 0.023865589344711436
k = 6   t = 7.448167562484741   F( $\beta_k$ ) = 424.9954901936109   || $\nabla F(\beta_k)$ ||_2= 0.00640751543983564   y_k = 0.04027129289545917
Algorithm ran for 7 iterations. Converged: True Training time: 7.448390245437622
Saving trained  $\beta$  in beta_large_gd
Reading test data from data/large.test
Reading  $\beta$  from beta_large_gd
Computing MSE on data data/large.test
MSE is: 438.2794383647582

```

## Question 8 :

### Part A : Function for solveLin()

```

def solveLin(z,K):
    """ Solve problem
        Minimize:  $z^T \beta$ 
        subject to:  $||\beta||_1 \leq K$ 
        The return value is the optimal  $\beta^*$ .
    """
    # Find coordinate  $k^*$  that maximizes  $|z_k|$ 
    k_star = np.argmax(np.abs(z))
    #Compute optimal beta
    beta_star = np.zeros_like(z)
    beta_star[k_star] = -K*np.sign(z[k_star])
    return beta_star

```

## Part B : The function for exact line search FW

```
def exactLineSearchFW(data,beta,s):
    """ Given data, a vector x, and a direction g, return
        | | |  $\gamma^*$  = argmin_{ $\gamma$  in [0,1]} F(data, (1- $\gamma$ ) $\beta$ + $\gamma$ s)

        The solution is found by first computing the coefficients of the quadratic
        polynomial
        | | |  $h(\gamma) = F(\text{data}, (1-\gamma)\beta + \gamma s) = a\gamma^2 + b\gamma + c$ 

        Inputs are:
        - data: an RDD containing pairs of the form (x,y)
        - beta: first interpolation vector  $\beta$ 
        - s: second interpolation vector s

        The return value is  $\gamma^*$ 

    """
    # here cof stores tuples cof=(a,b,c) obtained from hcoeff function
    cof = hcoeff(data, beta, s - beta)
    gamma = -cof[1] / (2 * cof[0])
    gamma = max(0, min(1, gamma))
    return gamma
```

## Part C:

```
def train_FW(data,beta_0, K,max_iter,eps):
    """ Use the Frank-Wolfe algorithm minimize  $F_0$  given by

        | | |  $F_0(\beta) = 1/n \sum_{(x,y) \text{ in data}} f(\beta;x,y)$ 

        Subject to:
        | | |  $\|\beta\|_1 \leq K$ 

        Inputs are:
        - data: an rdd containing pairs of the form (x,y)
        - beta_0: the starting vector  $\beta$ 
        - K: the bound K
        - max_iter: maximum number of iterations
        - eps: upper bound on the convergence criterion

        The function runs the Frank-Wolfe algorithm with a step-size found through
        exact line search. That is, it computes

        | | |  $s_k = \text{argmin}_{\{s: \|s\|_1 \leq K\}} s^T \nabla F_0(\beta_k)$ 
        | | |  $\beta_{k+1} = (1-\gamma_k)\beta_k + \gamma_k s_k$ 

        where the gain  $\gamma_k$  is given by

        | | |  $\gamma_k = \text{argmin}_{\gamma \text{ in } [0,1]} F_0((1-\gamma_k)\beta_k + \gamma_k s_k)$ 

        and terminates after max_iter iterations or when  $(\beta_k - s_k)^T \nabla F(\beta_k) < \epsilon$ .
```

```

    The function returns:
    -beta: the trained  $\beta$ ,
    -criterion: the condition  $(\beta_k - s_k)^T \nabla F(\beta_k)$ 
    -k: the number of iterations performed
"""
beta = beta_0
criterion = np.inf
k = 0
while k < max_iter and criterion > eps:
    grad = gradient(data, beta)
    s = solveLin(grad, K)
    gamma = exactLineSearchFW(data, beta, s)
    beta_next = (1 - gamma) * beta + gamma * s
    criterion = np.dot(beta - s, grad)
    beta = beta_next
    k += 1
return beta, criterion, k

```

**Question 9 :**

**Part A :** The following is the table of lambda with various values :

Lambda Value	MSE
0.125	187.4876923
0.25	180.9585623
0.5	184.0926789
1.0	207.4430064
2.0	251.8238198
4.0	304.6125705
8.0	350.7712854

The vector  $\beta$  for the smallest mse for  $\lambda=0.25$  is

7.021374726806946,5.39530071697655,3.8687673552936532,7.992597130017712,9.04928743708186  
9,1.6732785745748397,1.026298936238485,-0.8204970313544742,-  
0.18889666571174063,0.04677936033184126,-1.3354586513286817,-0.08615551747222905,-  
0.4373865737555297,0.35018289152900106,-0.22752018461561122,-0.6747107884658557,-  
0.6922443487619483,-1.1130039307621327,-0.6086923938363226,-  
0.9539084800499947,0.2090231391648798,0.5901512726555541,-  
0.21683279507970948,0.5248652600462588,-0.21167886242530973,0.06495615683450194,-  
0.6022594810140172

**Part B:**

K value	MSE
1.0	408.71093
5.0	361.7065
10.0	303.660
20.0	229.5094
30.0	188.8182
40.0	180.3845
50.0	184.1069

The vector  $\beta$  for the smallest MSE for  $k=40.0$  is

7.195908105527798,5.251560791266931,3.9150181134985504,9.01336708401366,10.0909500276657  
25,0.8477219404948313,0.1630402404570907,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,-  
0.5528261548856925,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0

**Part C :**

The difference between the two optimal solution is that in **b** there are more number of zeros present as compared to **a** . The reason can be due to following reasons such as convergence , usage of exact line search gradient and forward algorithm and Frank-Wolfe algorithm